

# boston-housing-data-analysis-day2

July 9, 2023

## 0.1 Small Project

Build a linear regression model to predict housing prices based on a given dataset.

- **Step 1:** Load and explore the dataset, including visualizing the features and target variable.
- **Step 2:** Split the dataset into training and testing sets.
- **Step 3:** Preprocess the data by handling missing values and performing feature scaling.
- **Step 4:** Train a linear regression model on the training data.
- **Step 5:** Evaluate the model's performance on the testing data using metrics such as mean squared error (MSE) or R-squared.

**crim:** Per capita crime rate by town

**zn:** Proportion of residential land zoned for lots over 25,000 sq.ft.

**indus:** Proportion of non-retail business acres per town.

**chas:** Charles River dummy variable (= 1 if tract bounds river; 0 otherwise).

**nox:** Nitrogen oxides concentration (parts per 10 million).

**rm:** Average number of rooms per dwelling.

**age:** Proportion of owner-occupied units built prior to 1940.

**dis:** Weighted mean of distances to five Boston employment centres.

**rad:** Index of accessibility to radial highways.

**tax:** Full-value property-tax rate per \$10,000.

**ptratio:** Pupil-teacher ratio by town.

**black:**  $1000(\text{Bk} - 0.63)^2$  where Bk is the proportion of blacks by town.

**lstat:** Lower status of the population (percent).

**medv:** Median value of owner-occupied homes in \$1000s.

## 0.2 Load Data

```
[7]: import pandas as pd
import matplotlib.pyplot as plt
df = pd.read_csv("train.csv")
df
```

```
[7]:
```

	ID	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	\
0	1	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	
1	2	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	
2	4	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	
3	5	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	
4	7	0.08829	12.5	7.87	0	0.524	6.012	66.6	5.5605	5	311	
..	...	...	...	...	...	...	...	...	...	...	...	...
328	500	0.17783	0.0	9.69	0	0.585	5.569	73.5	2.3999	6	391	
329	502	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	
330	503	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	
331	504	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	
332	506	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	

	ptratio	black	lstat	medv
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	18.7	394.63	2.94	33.4
3	18.7	396.90	5.33	36.2
4	15.2	395.60	12.43	22.9
..	...	...	...	...
328	19.2	395.77	15.10	17.5
329	21.0	391.99	9.67	22.4
330	21.0	396.90	9.08	20.6
331	21.0	396.90	5.64	23.9
332	21.0	396.90	7.88	11.9

[333 rows x 15 columns]

```
[8]: df.head(10)
```

```
[8]:
```

	ID	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	\
0	1	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	
1	2	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	
2	4	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	
3	5	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	
4	7	0.08829	12.5	7.87	0	0.524	6.012	66.6	5.5605	5	311	
5	11	0.22489	12.5	7.87	0	0.524	6.377	94.3	6.3467	5	311	
6	12	0.11747	12.5	7.87	0	0.524	6.009	82.9	6.2267	5	311	
7	13	0.09378	12.5	7.87	0	0.524	5.889	39.0	5.4509	5	311	
8	14	0.62976	0.0	8.14	0	0.538	5.949	61.8	4.7075	4	307	
9	15	0.63796	0.0	8.14	0	0.538	6.096	84.5	4.4619	4	307	

	ptratio	black	lstat	medv
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	18.7	394.63	2.94	33.4
3	18.7	396.90	5.33	36.2

4	15.2	395.60	12.43	22.9
5	15.2	392.52	20.45	15.0
6	15.2	396.90	13.27	18.9
7	15.2	390.50	15.71	21.7
8	21.0	396.90	8.26	20.4
9	21.0	380.02	10.26	18.2

### 0.3 Visualize Data using Scatter Charts



```
[137]: # df.scatter()
# plt.show()
# plt.tight_layout()

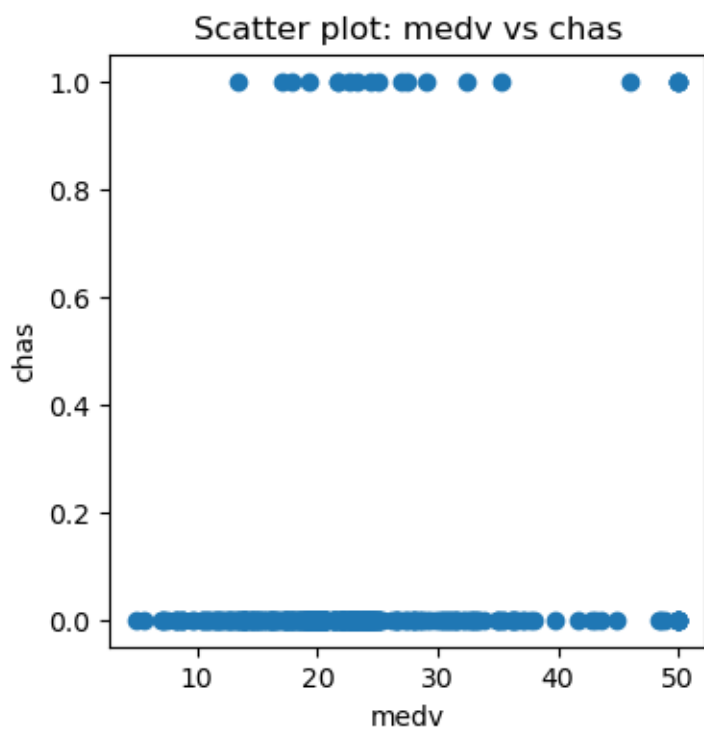

# plt.scatter(x=df['crim'], y=df['medv'])
# plt.xlabel("Crim")
# plt.ylabel("Medv")



for column in df.columns[1:]:
    plt.figure(figsize=(4,4))
    plt.scatter(df['medv'], df[column], cmap='Spectral', alpha=1)
    plt.xlabel('medv')
    plt.ylabel(column)
    plt.title(f'Scatter plot: medv vs {column}')
    plt.show()

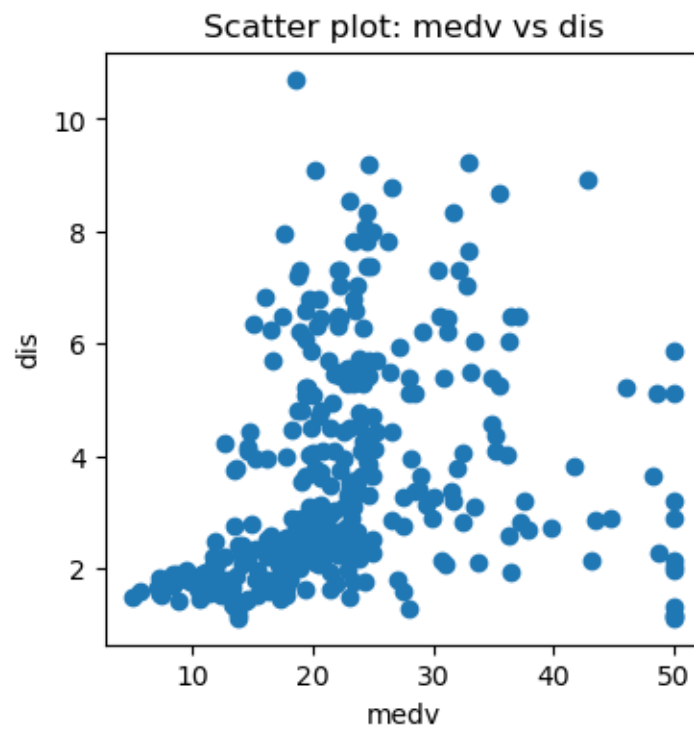

# for col in df.columns:
#     plt.scatter(x=df['medv'], y=col)
```



C:\Users\Hp\AppData\Local\Temp\ipykernel\_15644\105836776.py:11: UserWarning: No data for colormapping provided via 'c'. Parameters 'cmap' will be ignored  
plt.scatter(df['medv'], df[column], cmap='Spectral', alpha=1)













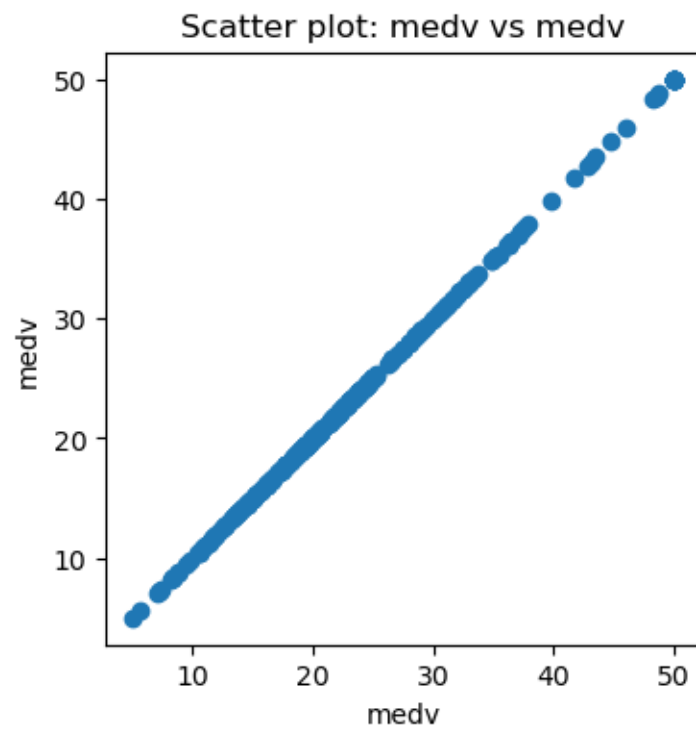


Scatter plot: medv vs ptratio



Scatter plot: medv vs black





#### 0.4 Fill Empty Data

```
[111]: df.medv.fillna(df.medv.mean(), inplace=True)
```

df

[111]:	ID	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	\
0	1	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	
1	2	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	
2	4	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	
3	5	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	
4	7	0.08829	12.5	7.87	0	0.524	6.012	66.6	5.5605	5	311	
..	...	...	...	...	...	...	...	...	...	...	...	
328	500	0.17783	0.0	9.69	0	0.585	5.569	73.5	2.3999	6	391	
329	502	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	
330	503	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	
331	504	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	
332	506	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	

	ptratio	black	lstat	medv
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	18.7	394.63	2.94	33.4
3	18.7	396.90	5.33	36.2
4	15.2	395.60	12.43	22.9
..	...	...	...	...
328	19.2	395.77	15.10	17.5
329	21.0	391.99	9.67	22.4
330	21.0	396.90	9.08	20.6
331	21.0	396.90	5.64	23.9
332	21.0	396.90	7.88	11.9

```
[333 rows x 15 columns]
```

### 0.5 Separate Training and Testing Data

```
[118]: from sklearn.model_selection import train_test_split

# Split arrays or matrices into random train and test subsets
x = df.drop('medv', axis=1)
y = df['medv']

X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
    ↪ random_state=25)

y_train
```

```
[118]: 204    33.4
        40    16.0
        23    34.9
        100   17.8
        126   39.8
        ...
        255    7.4
        317   23.7
        143   22.6
        318   21.8
        132   34.9
Name: medv, Length: 266, dtype: float64
```

## 0.6 Feature Scaling

```
[120]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

X_train_scaled
```

```
[120]: array([[ 0.39386662, -0.42470702,  0.90238868, ..., -0.02218378,
                0.43006366, -0.83126475],
               [-1.29151012, -0.41236258,  0.56749452, ...,  0.59645738,
                0.21981071,  0.25378715],
               [-1.43597098, -0.42999823,  2.66058302, ..., -0.06977156,
                0.41576378, -1.44254242],
               ...,
               [-0.30092134, -0.41686853, -0.47904973, ...,  0.07299179,
                0.43006366, -0.23223986],
               [ 1.61146533, -0.07442652, -0.47904973, ...,  0.83439629,
                0.38559996, -0.29350377],
               [-0.41098676, -0.42360001,  1.40472992, ..., -1.54499279,
                0.43006366, -0.97829813]])
```

## 0.7 Linear Regression and Visualizing Output

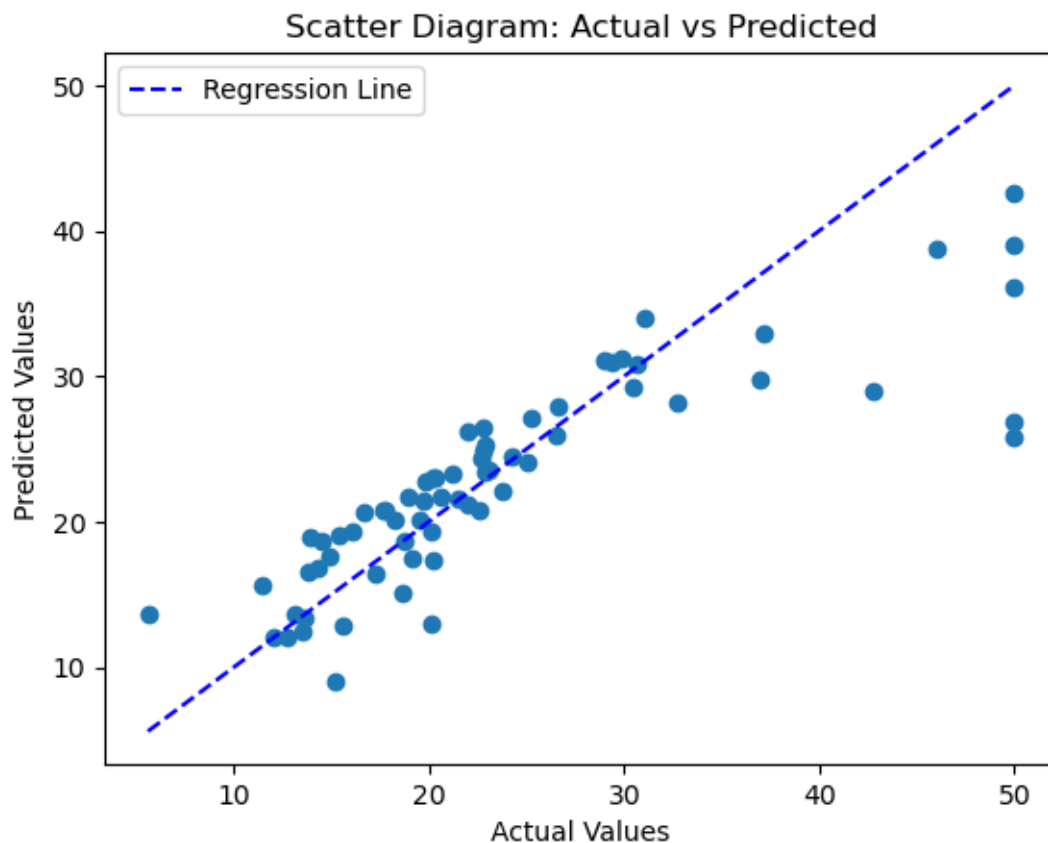
```
[128]: from sklearn.linear_model import LinearRegression

model = LinearRegression()
model.fit(X_train_scaled, y_train)
```

```
[128]: LinearRegression()
```

```
[133]: y_pred = model.predict(X_test_scaled)

# Plot the scatter diagram
plt.scatter(y_test, y_pred)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'b--',
         label='Regression Line')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('Scatter Diagram: Actual vs Predicted')
plt.legend()
plt.show()
```



## 0.8 Mean Square and R-squared

```
[136]: # Mean Square and R-squared
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
import math
```

```
mean_sqrt = math.sqrt(mean_squared_error(y_test, y_pred))
r_sqrt = r2_score(y_test,y_test)

print("Mean Squared {}".format(mean_sqrt))
print("R-Squared {}".format(r_sqrt))
```

```
Mean Squared 5.8183976341242385
R-Squared 1.0
```

```
[ ]:
```