

# Hypothesis\_KMeans

July 19, 2023

## 1 Task 1: Hypothesis testing and confidence intervals

Conduct hypothesis tests and calculate confidence intervals for the Heart Disease UCI dataset.

### 1.1 Load Packages

```
[1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

### 1.2 Read CSV Dataset

```
[2]: df = pd.read_csv("heart_disease_uci.csv")
df.head()
```

```
[2]:   id  age  sex  dataset      cp  trestbps  chol  fbs  \
0   1   63  Male  Cleveland  typical angina    145.0  233.0  True
1   2   67  Male  Cleveland  asymptomatic    160.0  286.0  False
2   3   67  Male  Cleveland  asymptomatic    120.0  229.0  False
3   4   37  Male  Cleveland  non-anginal    130.0  250.0  False
4   5   41  Female  Cleveland  atypical angina    130.0  204.0  False
```

```
      restecg  thalch  exang  oldpeak      slope  ca  \
0  lv hypertrophy  150.0  False      2.3  downsloping  0.0
1  lv hypertrophy  108.0   True      1.5      flat  3.0
2  lv hypertrophy  129.0   True      2.6      flat  2.0
3      normal  187.0  False      3.5  downsloping  0.0
4  lv hypertrophy  172.0  False      1.4  upsloping  0.0
```

```
      thal  num
0  fixed defect    0
1    normal     2
2  reversable defect  1
3    normal     0
4    normal     0
```

### 1.3 Data Wrangling/Cleaning

```
[120]: # drop empty stages (num) and cholesterol levels as we need both to be filled
df.dropna(subset=['chol', 'num'], inplace=True)
# df.num.isna().sum()
```

```
[121]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 890 entries, 0 to 919
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         890 non-null    int64
 1   sex         890 non-null    object
 2   dataset     890 non-null    object
 3   cp          890 non-null    object
 4   trestbps    834 non-null    float64
 5   chol        890 non-null    float64
 6   fbs         800 non-null    object
 7   restecg     888 non-null    object
 8   thalch      838 non-null    float64
 9   exang       838 non-null    object
10   oldpeak     831 non-null    float64
11   slope       603 non-null    object
12   ca          308 non-null    float64
13   thal        432 non-null    object
14   num         890 non-null    int64
dtypes: float64(5), int64(2), object(8)
memory usage: 143.5+ KB
```

```
[5]: df.dataset.value_counts()
```

```
[5]: Cleveland      304
Hungary           270
VA Long Beach     193
Switzerland       123
Name: dataset, dtype: int64
```

```
[6]: df.num.value_counts()
```

```
[6]: 0    392
1    258
2    107
3    106
4     27
Name: num, dtype: int64
```

```
[7]: df.isna().sum()
```

```
[7]: id          0
     age         0
     sex         0
     dataset     0
     cp          0
     trestbps    56
     chol        0
     fbs        90
     restecg     2
     thalch      52
     exang       52
     oldpeak     59
     slope      287
     ca         582
     thal       458
     num         0
     dtype: int64
```

```
[8]: df.describe()
```

```
[8]:
```

	id	age	trestbps	chol	thalch	oldpeak	\
count	890.000000	890.000000	834.000000	890.000000	838.000000	831.000000	
mean	458.016854	53.580899	132.089928	199.130337	137.539379	0.889290	
std	267.339571	9.389502	19.077093	110.780810	25.989709	1.095398	
min	1.000000	28.000000	0.000000	0.000000	60.000000	-2.600000	
25%	223.250000	47.000000	120.000000	175.000000	120.000000	0.000000	
50%	461.500000	54.000000	130.000000	223.000000	140.000000	0.500000	
75%	690.750000	60.000000	140.000000	268.000000	157.000000	1.500000	
max	920.000000	77.000000	200.000000	603.000000	202.000000	6.200000	

	ca	num
count	308.000000	890.000000
mean	0.678571	1.008989
std	0.936378	1.145210
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	1.000000
75%	1.000000	2.000000
max	3.000000	4.000000

## 1.4 Find Co-Relations

```
[9]: df.corr()
```

```
[9]:
```

	id	age	trestbps	chol	thalch	oldpeak	\
id	1.000000	0.230532	0.053213	-0.376936	-0.470599	0.043070	
age	0.230532	1.000000	0.253467	-0.086234	-0.360682	0.252412	
trestbps	0.053213	0.253467	1.000000	0.092853	-0.114104	0.170562	
chol	-0.376936	-0.086234	0.092853	1.000000	0.236121	0.047734	
thalch	-0.470599	-0.360682	-0.114104	0.236121	1.000000	-0.151671	
oldpeak	0.043070	0.252412	0.170562	0.047734	-0.151671	1.000000	
ca	0.068357	0.372018	0.094925	0.051606	-0.265275	0.280301	
num	0.267503	0.329766	0.128628	-0.231547	-0.371710	0.434298	


	ca	num
id	0.068357	0.267503
age	0.372018	0.329766
trestbps	0.094925	0.128628
chol	0.051606	-0.231547
thalch	-0.265275	-0.371710
oldpeak	0.280301	0.434298
ca	1.000000	0.515338
num	0.515338	1.000000

## 1.5 Co-Relation Heatmap

```
[10]: correlations = df.corr()

sns.heatmap(correlations, annot=True, fmt=".2f")
```

```
[10]: <Axes: >
```



```
[11]: # dropping unnecessary
df.drop(['id'], axis='columns', inplace=True)
```

```
[12]: df
```

```
[12]:
```

	age	sex	dataset	cp	trestbps	chol	fbs	\
0	63	Male	Cleveland	typical angina	145.0	233.0	True	
1	67	Male	Cleveland	asymptomatic	160.0	286.0	False	
2	67	Male	Cleveland	asymptomatic	120.0	229.0	False	
3	37	Male	Cleveland	non-anginal	130.0	250.0	False	
4	41	Female	Cleveland	atypical angina	130.0	204.0	False	
..	...	...	...	...	...	...		
915	54	Female	VA Long Beach	asymptomatic	127.0	333.0	True	
916	62	Male	VA Long Beach	typical angina	NaN	139.0	False	
917	55	Male	VA Long Beach	asymptomatic	122.0	223.0	True	
918	58	Male	VA Long Beach	asymptomatic	NaN	385.0	True	
919	62	Male	VA Long Beach	atypical angina	120.0	254.0	False	

	restecg	thalch	exang	oldpeak	slope	ca	\
0	lv hypertrophy	150.0	False	2.3	downsloping	0.0	
1	lv hypertrophy	108.0	True	1.5	flat	3.0	
2	lv hypertrophy	129.0	True	2.6	flat	2.0	

3	normal	187.0	False	3.5	downsloping	0.0
4	lv hypertrophy	172.0	False	1.4	upsloping	0.0
..	...	...	...	...	...	...
915	st-t abnormality	154.0	False	0.0	NaN	NaN
916	st-t abnormality	NaN	NaN	NaN	NaN	NaN
917	st-t abnormality	100.0	False	0.0	NaN	NaN
918	lv hypertrophy	NaN	NaN	NaN	NaN	NaN
919	lv hypertrophy	93.0	True	0.0	NaN	NaN

	thal	num
0	fixed defect	0
1	normal	2
2	reversable defect	1
3	normal	0
4	normal	0
..	...	...
915	NaN	1
916	NaN	0
917	fixed defect	2
918	NaN	0
919	NaN	1


[890 rows x 15 columns]

## 1.6 Visualize Different Details

```
[13]: pd.crosstab(df.num, df.sex)
```


```
[13]: sex  Female  Male
num
0      137    255
1       30    228
2       10     97
3        8     98
4         2     25
```

```
[14]: pd.crosstab(df.num, df.sex).plot(kind="bar", color=['salmon', 'lightblue'])
plt.xlabel("Stages (0=No Disease)")
plt.ylabel("Amount of People")
plt.show()
```



```
[15]: # Compare Age vs max Heart Rate(thalch)
plt.scatter(df.age[df.num==0], df.thalch[df.num==0], c='green')
plt.scatter(df.age[df.num==1], df.thalch[df.num==1], c='blue')
plt.scatter(df.age[df.num==2], df.thalch[df.num==2], c='salmon')
plt.scatter(df.age[df.num==3], df.thalch[df.num==3], c='lightblue')
plt.scatter(df.age[df.num==4], df.thalch[df.num==4], c='red')

plt.xlabel("Age")
plt.ylabel("Max Heart Rate")
plt.legend(['No Disease', 'Stage 1', 'Stage 2', 'Stage 3', 'Stage 4'])
plt.show()
```




```
[16]: # Heart Disease per chest pain
pd.crosstab(df.cp, df.num)
```

```
[16]: num          0    1    2    3    4
cp
asymptomatic    102  190   87   82   23
atypical angina  140   19    2    3    0
non-anginal     124   37   14   18    3
typical angina   26   12    4    3    1
```

```
[17]: pd.crosstab(df.cp, df.num).plot(kind="bar", color=['green', 'blue', 'salmon', 'lightblue', 'red'])
plt.ylabel("Amount")
plt.show()
```





```
[18]: df
```

```
[18]:
```

	age	sex	dataset	cp	trestbps	chol	fbs	\
0	63	Male	Cleveland	typical angina	145.0	233.0	True	
1	67	Male	Cleveland	asymptomatic	160.0	286.0	False	
2	67	Male	Cleveland	asymptomatic	120.0	229.0	False	
3	37	Male	Cleveland	non-anginal	130.0	250.0	False	
4	41	Female	Cleveland	atypical angina	130.0	204.0	False	
..	..	..	..	..	..	..	..	
915	54	Female	VA Long Beach	asymptomatic	127.0	333.0	True	
916	62	Male	VA Long Beach	typical angina	NaN	139.0	False	
917	55	Male	VA Long Beach	asymptomatic	122.0	223.0	True	
918	58	Male	VA Long Beach	asymptomatic	NaN	385.0	True	
919	62	Male	VA Long Beach	atypical angina	120.0	254.0	False	

```

restecg thalch exang oldpeak      slope  ca  \

```

0	lv hypertrophy	150.0	False	2.3	downsloping	0.0
1	lv hypertrophy	108.0	True	1.5	flat	3.0
2	lv hypertrophy	129.0	True	2.6	flat	2.0
3	normal	187.0	False	3.5	downsloping	0.0
4	lv hypertrophy	172.0	False	1.4	upsloping	0.0
..	...	...	...	...	...	...
915	st-t abnormality	154.0	False	0.0	NaN	NaN
916	st-t abnormality	NaN	NaN	NaN	NaN	NaN
917	st-t abnormality	100.0	False	0.0	NaN	NaN
918	lv hypertrophy	NaN	NaN	NaN	NaN	NaN
919	lv hypertrophy	93.0	True	0.0	NaN	NaN

	thal	num
0	fixed defect	0
1	normal	2
2	reversable defect	1
3	normal	0
4	normal	0
..	...	...
915	NaN	1
916	NaN	0
917	fixed defect	2
918	NaN	0
919	NaN	1

[890 rows x 15 columns]

## 1.7 Finding Hypothesis

### 1.8 Null Hypothesis (H0):

There is no significant difference in cholesterol levels between patients with and without heart disease.

### 1.9 Alternative Hypothesis (H1):

There is a significant difference in cholesterol levels between patients with and without heart disease.

```
[79]: age_heart_disease = df.groupby("num")["age"]

fig, (axis1, axis2, axis3, axis4, axis5) = plt.subplots(1, 5, figsize=(20, 5))


ax = sns.histplot(age_heart_disease.get_group(0), kde=True, color='green',
    ↪ax=axis1)
ax.set(xlabel="Age with No Disease")
ax = sns.histplot(age_heart_disease.get_group(1), kde=True, color='salmon',
    ↪ax=axis2)
ax.set(xlabel="Age with Stage 1")
```

```

ax = sns.histplot(age_heart_disease.get_group(2), kde=True, color='#F70000',
    ↪ax=axis3)
ax.set(xlabel="Age with Stage 2")
ax = sns.histplot(age_heart_disease.get_group(3), kde=True, color='#B10000',
    ↪ax=axis4)
ax.set(xlabel="Age with Stage 3")
ax = sns.histplot(age_heart_disease.get_group(4), kde=True, color='#960000',
    ↪ax=axis5)
ax.set(xlabel="Age with Stage 4")

plt.show()

```



```

[80]: # Hypothesis Testing
ages_mean = df.groupby("num")['age'].mean()
ages_std = df.groupby("num")['age'].std()

stages_df = pd.DataFrame({'Stage':[0,1,2,3,4], 'Age Mean': ages_mean.values,
    ↪'Age Std': ages_std.values,
                                'Sample Size':[ len(df.age[df.num==0]), len(df.age[df.
    ↪num==1]), len(df.age[df.num==2]),
                                                len(df.age[df.num==3]), len(df.age[df.
    ↪num==4]) ] })
stages_df

```

```

[80]:   Stage  Age Mean  Age Std  Sample Size
0      0  50.704082  9.507525           392
1      1  53.511628  8.755757           258
2      2  57.336449  7.650310           107
3      3  59.320755  7.952894           106
4      4  58.592593  7.747254            27

```

```

[103]: # Population Density Function of Age

sns.histplot(age_heart_disease.get_group(0), label='No Heart Disease', kde=True)
sns.histplot(pd.concat([age_heart_disease.get_group(1),


```

```

        age_heart_disease.get_group(2),
        age_heart_disease.get_group(3),
        age_heart_disease.get_group(4)]), label='Yes Heart_
↪Disease', kde=True)
plt.legend()
plt.show()

# we can observe that the distribution of the age of the person who doesn't_
↪have heart disease
# shifted downward and to the left of those who have heart disease.

```



Since we want to compare the cholesterol levels (numerical variable) between two groups (patients with and without heart disease), we will use an independent t-test.

```

[104]: from scipy.stats import ttest_ind

noheart_disease = df[df.num == 0]['chol']
heart_disease = df[df.num > 0]['chol']

# perform independent t-test
t_stat, p_value = ttest_ind(heart_disease, noheart_disease)

```

```

# Print the results
print("t-statistic:", t_stat)
print("p-value:", p_value)

alpha = 0.05

if p_value < alpha:
    print("Reject Null Hypothesis. There is a significant difference in_
    ↪cholesterol levels.")
else:
    print("Fail to Reject Null Hypothesis. There is no significant difference_
    ↪in cholesterol levels.")

```

t-statistic: -7.061510014417096

p-value: 3.323201042049721e-12

Reject Null Hypothesis. There is a significant difference in cholesterol levels.

## 1.10 Confidence Intervals

```

[119]: # Confidence Interval
# confidence intervals for the 'chol' (cholesterol levels) and 'thalach'_
    ↪(maximum heart rate achieved) variables
import math

columns = df[['thalch', 'chol']]
columns_mean = columns.mean()
columns_std = columns.std()

columns_standard_error = columns_std/len(columns)
columns_margin_error = columns_standard_error/2

confidence_level = 0.95

upper_bound = columns_mean + confidence_level * (columns_std / math.sqrt(10))
lower_bound = columns_mean - confidence_level * (columns_std / math.sqrt(10))

confidence_intervals = pd.DataFrame({
    'Sample Size': columns.count(),
    'Sample Mean': columns_mean,
    'Standard Error': columns_standard_error,
    'Margin of Error': columns_margin_error,
    'Lower Bound (95% CI)': lower_bound,
    'Upper Bound (95% CI)': upper_bound
})

confidence_intervals

```

```
[119]:
```

	Sample Size	Sample Mean	Standard Error	Margin of Error \
thalch	838	137.539379	0.029202	0.014601
chol	890	199.130337	0.124473	0.062236

	Lower Bound (95% CI)	Upper Bound (95% CI)
thalch	129.731645	145.347114
chol	165.849967	232.410707


```
[ ]:
```

## 2 Task 2: K-Means: Document the results for K=2, 3, 5, 10, 15 and 20. Comment on the results.

### 2.1 Load Image

```
[167]: import cv2
import numpy as np

faceImage = cv2.imread("face.jpg")
faceImage = cv2.cvtColor(faceImage, cv2.COLOR_BGR2RGB)
plt.imshow(faceImage)
plt.show()
```



## 2.2 Convert Image to Numerical Pixels Values

```
[168]: # Reshape the image to a 2D array of pixels
pixel_values = faceImage.reshape((-1, 3))
# pixel_values = np.float32(pixel_values)

pixel_values
print(pixel_values.shape)
```

(44000, 3)

## 2.3 Clustering (Getting the Cluster Centers) and Image Segmentation

```
[212]: from sklearn.cluster import KMeans
import time

plt.figure(figsize=(12, 8))

#clusters values
K_values = [2, 3, 5, 10, 15, 20]

for i, K in enumerate(K_values):
    time_start = time.time()

    # CLUSTERING
    kmeans = KMeans(n_clusters = K, n_init='auto')
    kmeans.fit(pixel_values)

    # Get the cluster centers and labels
    cluster_centers = kmeans.cluster_centers_
    labels = kmeans.labels_

    # SEGMENTATION

    # Create a new image with the same shape as the original image
    segmented_image = np.zeros_like(pixel_values)

    # Replace each pixel value with the corresponding cluster center
    for j in range(len(cluster_centers)):
        segmented_image[labels == j] = cluster_centers[j]

    # Reshape the segmented_image back to the original image shape
    segmented_image = segmented_image.reshape(faceImage.shape)

    time_end = time.time()
    computational_time = time_end - time_start
```


```

# segmented_image
plt.subplot(1, len(K_values), i + 1) # Use i + 1 for the subplot index
plt.imshow(segmented_image)

plt.axis('off')
plt.text(10, 250, f'Total time {round(computational_time, 2)}', fontsize=12)
plt.title(f'K={K}')

plt.tight_layout()
plt.show()

```



## 2.4 Observations

**Number Of Segments** As you increase the value of K, the number of segments or clusters in the segmented image increases. Smaller K values might group similar regions together, resulting in larger segments, whereas larger K values tend to create more detailed and smaller segments.

**Detail Level** Higher K values generally produce more detailed segmented images with finer distinctions between different regions or objects in the image. Lower K values, on the other hand, might merge similar regions, leading to a loss of fine details.

**Computational Complexity** The computational cost of k-means clustering increases with higher K values. Larger K values require more iterations and may be computationally expensive, especially for large images.

## 3 Task 3 Read and make a report

Link <https://aws.amazon.com/getting-started/hands-on/build-train-deploy-machine-learning-model-sagemaker/>

### 3.1 Here are the main steps covered in the article:

#### 3.1.1 Create a SageMaker notebook instance

You start by creating a SageMaker notebook instance where you can work on your data and ML model. This instance is used to download and process the data.



### 3.1.2 Prepare the data

The tutorial uses the Bank Marketing Data Set, which contains customer demographics and responses to marketing events. The data is pre-labeled to identify whether a customer enrolls for a product offered by the bank.

### 3.1.3 Train the model

After preparing the data, you use gradient-based optimization to train the XGBoost model on the data. The trained model will predict whether a customer will enroll for a certificate of deposit (CD).

### 3.1.4 Deploy the model

Once the model is trained, you deploy it to a SageMaker endpoint, making it accessible for predictions.

### 3.1.5 Evaluate model performance

The tutorial evaluates the model's performance using a confusion matrix, comparing actual and predicted values.

### 3.1.6 Clean up

Finally, you clean up by deleting the endpoint and S3 bucket used in the tutorial.

The tutorial also provides code snippets and explanations for each step. The cost is less than \$1, making it AWS Free Tier eligible.

[ ]: