

Customer Segmentation

August 20, 2023

```
[52]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[179]: data_size = 1000

data = {
    'customer_id': list(range(1000, 2000)),
    'age': np.random.randint(18, 100, data_size),
    'gender': np.random.choice(['Male', 'Female'], size=data_size),
    'marital_status': np.random.choice(['Married', 'Single', 'Divorced'],
    ↪size=data_size),
    'annual_income_in_usd': np.random.randint(25000, 150000, data_size),
    'total_purchases': np.random.randint(1, 50, data_size),
    'preferred_category': np.random.choice(['Electronics', 'Appliances',
    ↪'Clothing', 'Food', 'Health', 'Media'], size=data_size)
}
```

```
[180]: df = pd.DataFrame(data)
```

```
[181]: df.head()
```

```
[181]:
```

	customer_id	age	gender	marital_status	annual_income_in_usd	\
0	1000	40	Female	Married	106098	
1	1001	52	Male	Single	142620	
2	1002	37	Female	Divorced	107244	
3	1003	82	Male	Married	65595	
4	1004	60	Male	Divorced	79122	

	total_purchases	preferred_category
0	11	Food
1	19	Food
2	41	Media
3	15	Clothing
4	28	Electronics

```
[182]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customer_id           1000 non-null   int64
1   age                   1000 non-null   int32
2   gender                 1000 non-null   object
3   marital_status        1000 non-null   object
4   annual_income_in_usd  1000 non-null   int32
5   total_purchases       1000 non-null   int32
6   preferred_category    1000 non-null   object
dtypes: int32(3), int64(1), object(3)
memory usage: 43.1+ KB

```

```
[183]: df.describe()
```

```

[183]:      customer_id      age  annual_income_in_usd  total_purchases
count  1000.000000  1000.000000      1000.000000      1000.000000
mean    1499.500000    58.324000      86599.276000       25.112000
std      288.819436    23.744659      35019.782397       14.059414
min     1000.000000    18.000000      25022.000000        1.000000
25%     1249.750000    37.000000      56976.250000       13.000000
50%     1499.500000    57.500000      84772.000000       25.000000
75%     1749.250000    79.000000     115954.500000       37.000000
max     1999.000000    99.000000     149997.000000       49.000000

```

```
[184]: df.isna().sum()
```


```

[184]: customer_id      0
      age            0
      gender          0
      marital_status  0
      annual_income_in_usd  0
      total_purchases  0
      preferred_category  0
      dtype: int64

```


```
[185]: sns.histplot(df.age)
```

```
[185]: <Axes: xlabel='age', ylabel='Count'>
```




```
[186]: sns.boxplot(data=df, x='gender', y='annual_income_in_usd')
```

```
[186]: <Axes: xlabel='gender', ylabel='annual_income_in_usd'>
```



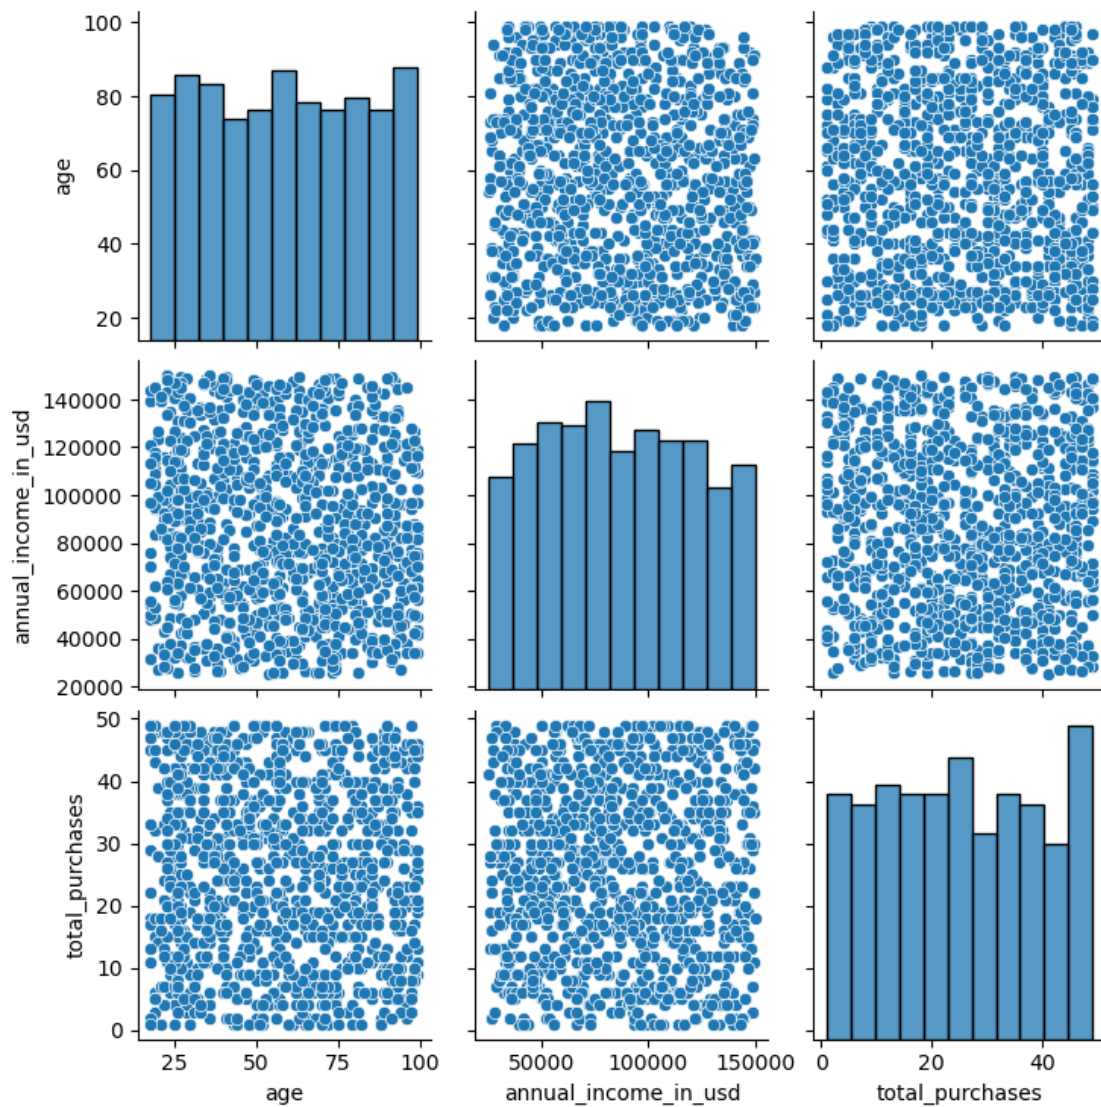
```
[187]: sns.countplot(data=df, x='marital_status')
```

```
[187]: <Axes: xlabel='marital_status', ylabel='count'>
```




```
[188]: sns.pairplot(df[['age', 'annual_income_in_usd', 'total_purchases']])
```

```
[188]: <seaborn.axisgrid.PairGrid at 0x2b0d1dfcb80>
```




```
[189]: sns.heatmap(df[['age', 'annual_income_in_usd', 'total_purchases']].corr(),
          ↳annot=True)
```

```
[189]: <Axes: >
```




```
[190]: sns.scatterplot(data=df[['age', 'annual_income_in_usd', 'total_purchases']],  
    ↪ x='age', y='annual_income_in_usd', hue='total_purchases')
```

```
[190]: <Axes: xlabel='age', ylabel='annual_income_in_usd'>
```



```
[191]: sns.jointplot(data=df, x='age', y='annual_income_in_usd', kind='hex')  
plt.show()
```

1 K Means

```
[192]: # features
X = df[['annual_income_in_usd', 'total_purchases']].values
X
```

```
[192]: array([[106098,    11],
              [142620,    19],
              [107244,    41],
              ...,
              [ 99362,    33],
              [ 66496,     3],
              [121458,    16]])
```

```
from sklearn.cluster import KMeans
```


```
# finding the right amount of clusters for kmeans using elbow method
wcss_list = []

for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, init='k-means++', random_state=42)
    kmeans.fit(X)
    wcss_list.append(kmeans.inertia_)

plt.plot(range(1, 11), wcss_list)
plt.show()
```

[illegible]

```
warnings.warn(
D:\apps\anaconda\files\lib\site-packages\sklearn\cluster\_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
```




```
[195]: # K means clustering and its training
kmeans = KMeans(n_clusters=3, init='k-means++', random_state=42)
y_predict = kmeans.fit_predict(X)
```

```
D:\apps\anaconda\files\lib\site-packages\sklearn\cluster\_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
```

```
[196]: # visualizing each cluster and their respective centers
plt.scatter(X[y_predict==0, 0], X[y_predict==0, 1], s=20, c='magenta',
            ↪label='Cluster 1')
plt.scatter(X[y_predict==1, 0], X[y_predict==1, 1], s=20, c='red',
            ↪label='Cluster 2')
plt.scatter(X[y_predict==2, 0], X[y_predict==2, 1], s=20, c='green',
            ↪label='Cluster 3')
```

```
plt.scatter(kmeans.cluster_centers[:, 0], kmeans.cluster_centers[:, 1], s = 100, c = 'blue', label = 'Cluster center')

plt.title('Clusters of customers')
plt.xlabel('Annual Income($)')
plt.ylabel('Purchases')
plt.legend(bbox_to_anchor=(0.9, 0., 0.5, 0.5))
plt.show()
```



2 Summary

In our analysis using K-means clustering, we aimed to uncover underlying patterns and groupings within our dataset based on the features we selected. The goal was to identify distinct clusters of customers that exhibit similar behavior in terms of annual income and purchases.

Upon running the K-means clustering algorithm and examining the resulting visualization, we observed that the three clusters we defined seem to exhibit nearly identical patterns. This outcome raises several interesting insights and considerations.

The output image is clearly showing the three different clusters with different colors. The clusters are formed between two parameters of the dataset; Annual income of customer and Purchases. We can change the colors and labels as per the requirement or choice. We can also observe some points from the above patterns, which are given below:

Similar Patterns Across Clusters: The fact that all three clusters display comparable graphs suggests that the algorithm faced challenges in distinguishing clear and separate patterns within the data using the chosen features. While we had hoped to find distinct customer segments based on annual income and purchases, the data does not appear to exhibit these expected separations.

Potential Explanations: There could be several reasons behind this outcome. One possibility is that the features we selected are not sufficient to effectively differentiate between the clusters. This might indicate that additional factors or features need to be considered to accurately capture the nuances in customer behavior.

In conclusion, the current K-means clustering analysis has given us a starting point for understanding customer behavior based on annual income and purchases. While the initial results did not yield distinct clusters, they have provided us with directions for further exploration and refinement of our analysis methodology.

[]:

[]: