

Customer Churn Prediction

August 22, 2023

1 Customer Churn Prediction for Telecommunications Company

Client Name: CommLink Telecom

Company Name: DataSense Solutions **Description:** CommLink Telecom, a telecommunications company, is facing high customer churn rates and wants to address the issue proactively. They have engaged DataSense Solutions to build a customer churn prediction model that can identify customers likely to churn in the near future. This will enable them to take targeted retention measures and improve customer retention rates.

Dataset:

CustomerID	Gender	Age	ServiceLength (months)	ContractType	MonthlyCharges (USD)	TotalCharges (USD)	Churn
1001	Male	42	24	Two-Year	85.00	2040.00	No
1002	Female	35	12	One-Year	79.50	942.50	Yes
1003	Male	62	48	Month-to-Month	94.20	4567.75	Yes
1004	Female	52	36	One-Year	78.25	2853.50	No
1005	Male	28	6	Month-to-Month	68.75	452.25	No
...

(Note: The dataset contains a total of 1000 customers with some churned (Yes) and others active (No).)

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

2 Generate Dataset

```
[2]: data_size = 1000
np.random.seed(42)

srv_length = np.random.randint(1, 50, data_size)
```

```

monthly_charges = np.random.uniform(20, 100, size=1000)

# more realistic distribution of churn
churn_probabilities = [0.4, 0.6]
churn = np.random.choice(['Yes', 'No'], size=data_size, p=churn_probabilities)

data = {
    'CustomerID': list(range(1001, 2001)),
    'Gender': np.random.choice(['Male', 'Female'], size=data_size),
    'Age': np.random.randint(18, 70, data_size),
    'ServiceLength (months)': srv_length,
    'ContractType': np.random.choice(['Two-Year', 'One-Year', 'Month-to-Month'], size=data_size),
    'MonthlyCharges (USD)': monthly_charges,
    'TotalCharges (USD)': srv_length * monthly_charges * 0.9,
    'Churn': churn
}

df = pd.DataFrame(data)

```

```
[3]: df.head()
```

```

[3]:   CustomerID  Gender  Age  ServiceLength (months)  ContractType \
0         1001   Male   66                39      Two-Year
1         1002  Female   27                29  Month-to-Month
2         1003   Male   45                15      Two-Year
3         1004   Male   38                43      Two-Year
4         1005   Male   67                 8      Two-Year

   MonthlyCharges (USD)  TotalCharges (USD)  Churn
0          75.756917      2659.067786    Yes
1          99.780443      2604.269557     No
2          91.728821      1238.339084    Yes
3          66.079873      2557.291102    Yes
4          93.391649       672.419873    Yes

```

```
[4]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CustomerID                          1000 non-null   int64
1   Gender                              1000 non-null   object
2   Age                                 1000 non-null   int32
3   ServiceLength (months)              1000 non-null   int32
4   ContractType                        1000 non-null   object

```

```

5   MonthlyCharges (USD)    1000 non-null   float64
6   TotalCharges (USD)      1000 non-null   float64
7   Churn                   1000 non-null   object
dtypes: float64(2), int32(2), int64(1), object(3)
memory usage: 54.8+ KB

```

```
[5]: df.describe()
```

```

[5]:      CustomerID      Age  ServiceLength (months)  MonthlyCharges (USD) \
count    1000.000000    1000.000000           1000.000000           1000.000000
mean     1500.500000     43.947000             25.379000             60.812403
std       288.819436     14.853984             14.122138             22.597415
min       1001.000000     18.000000              1.000000             20.019002
25%      1250.750000     32.000000             13.000000             41.524641
50%      1500.500000     44.000000             26.000000             62.128091
75%      1750.250000     57.000000             37.000000             79.855529
max       2000.000000     69.000000             49.000000             99.948280

      TotalCharges (USD)
count           1000.000000
mean           1374.978457
std             960.461013
min             18.074123
25%            564.412028
50%           1205.314557
75%           1944.868429
max            4295.379738

```

```
[6]: df.isna().sum()
```

```


[6]: CustomerID           0
     Gender              0
     Age                0
     ServiceLength (months)  0
     ContractType        0
     MonthlyCharges (USD)  0
     TotalCharges (USD)   0
     Churn               0
dtype: int64

```

3 Visualizationa

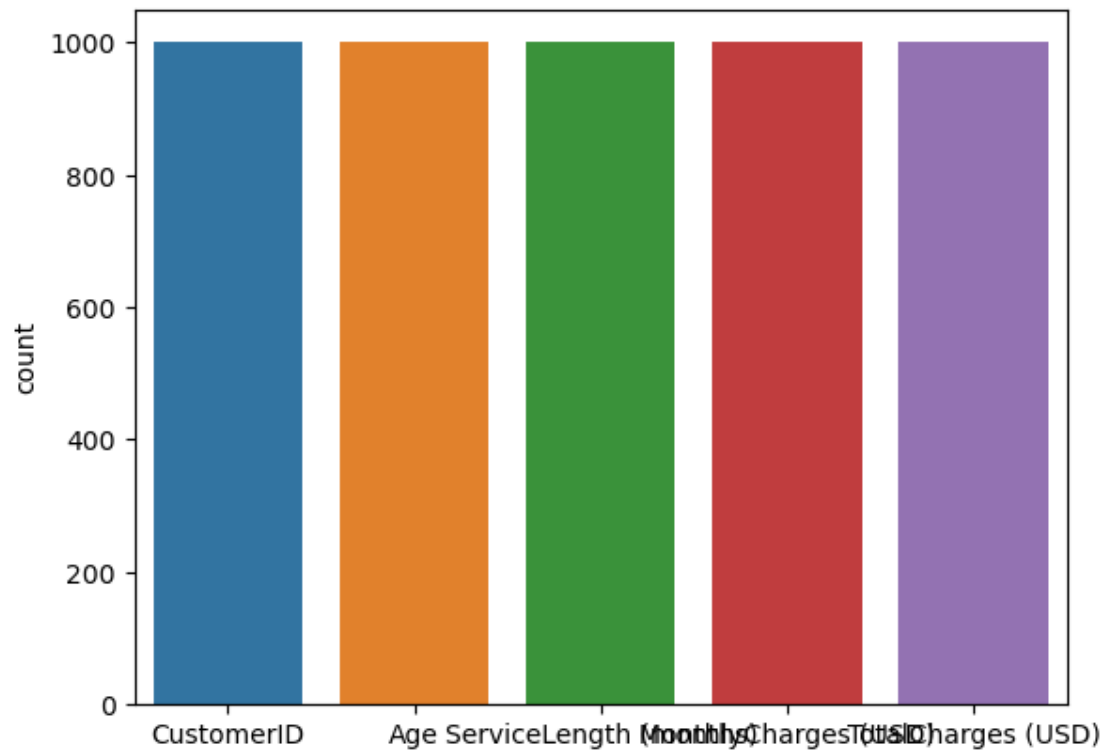
```
[7]: sns.countplot(data=df, x='Gender')
```

```
[7]: <Axes: xlabel='Gender', ylabel='count'>
```



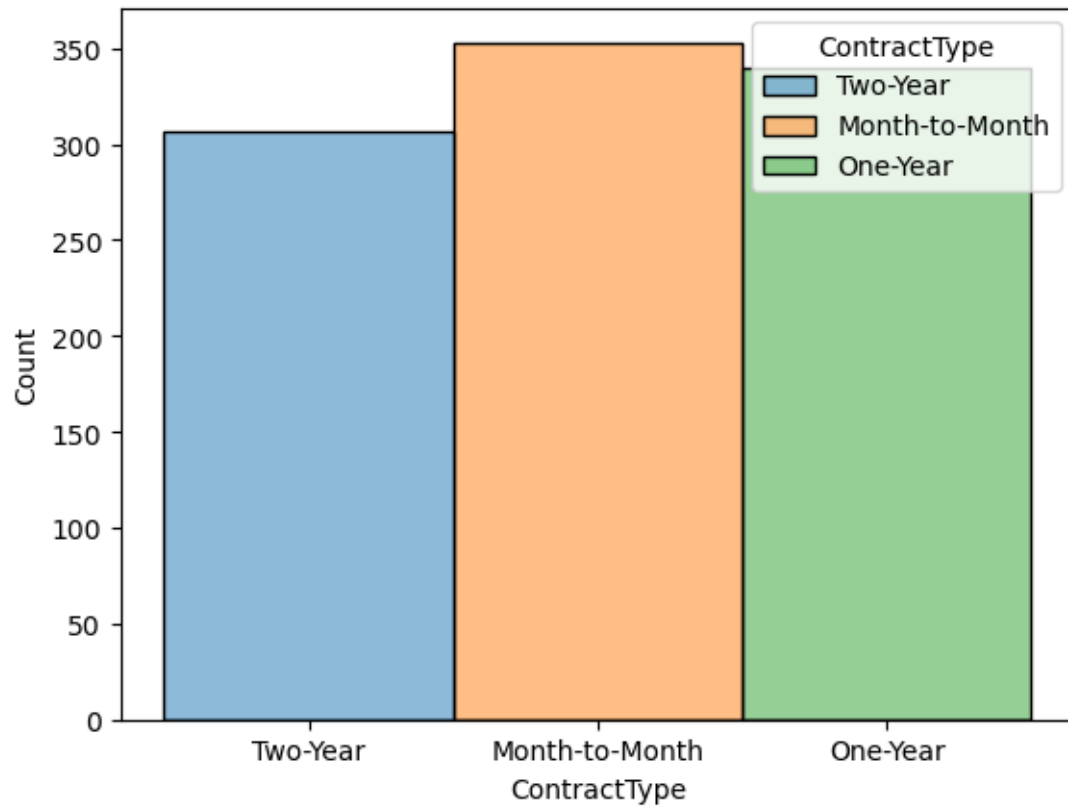
```
[8]: sns.countplot(data=df)
```

```
[8]: <Axes: ylabel='count'>
```




```
[9]: sns.histplot(data=df, x='ContractType', hue='ContractType')
```

```
[9]: <Axes: xlabel='ContractType', ylabel='Count'>
```



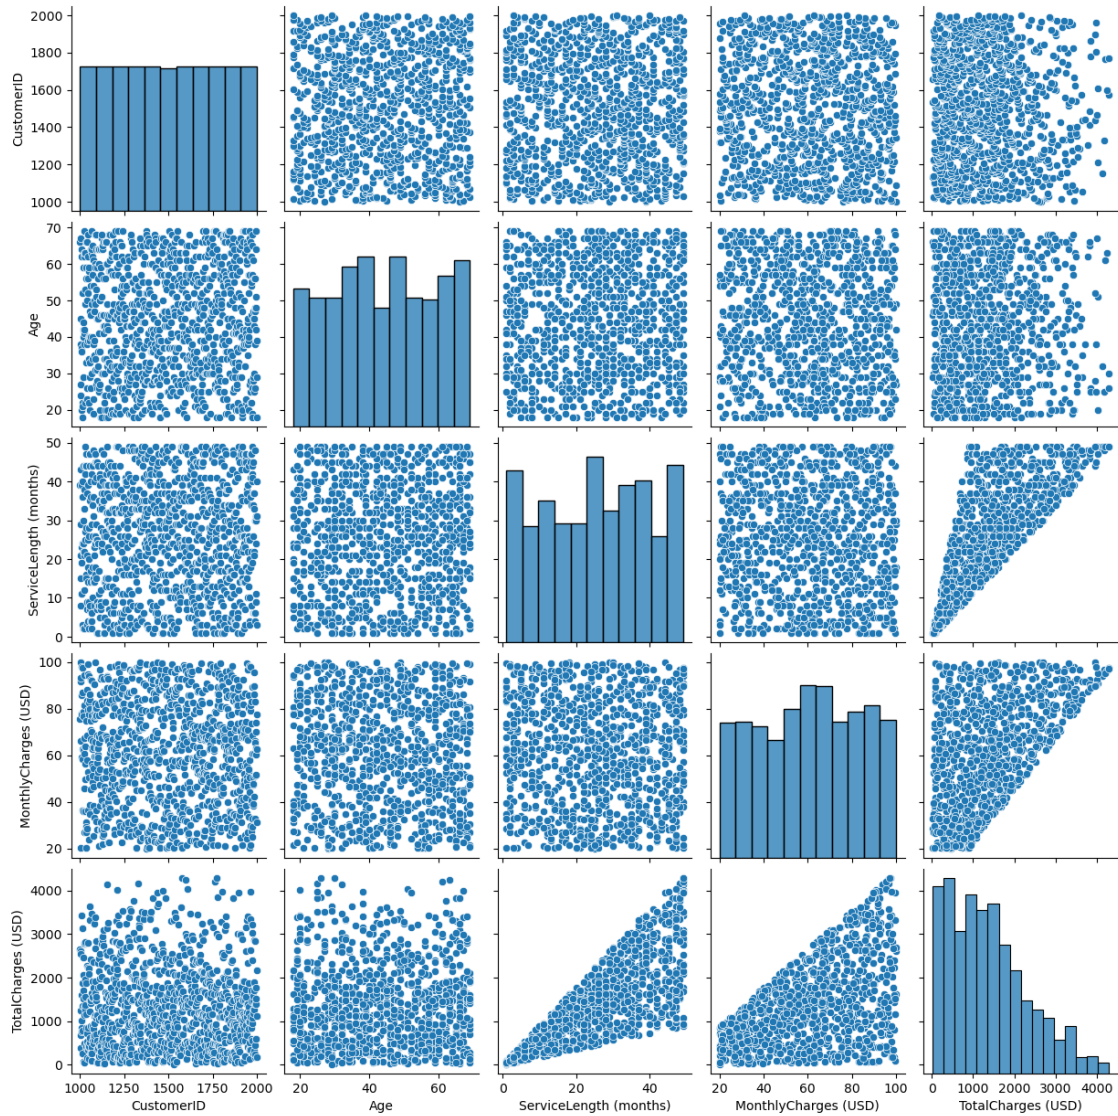
```
[10]: sns.scatterplot(data=df, x='Age', y='TotalCharges (USD)', hue='Age')
```

```
[10]: <Axes: xlabel='Age', ylabel='TotalCharges (USD)'>
```



```
[11]: sns.pairplot(df)
```

```
[11]: <seaborn.axisgrid.PairGrid at 0x12d4352f340>
```



4 Data Cleaning

[12]: df

```
[12]:
```

	CustomerID	Gender	Age	ServiceLength (months)	ContractType \
0	1001	Male	66	39	Two-Year
1	1002	Female	27	29	Month-to-Month
2	1003	Male	45	15	Two-Year
3	1004	Male	38	43	Two-Year
4	1005	Male	67	8	Two-Year
..
995	1996	Female	69	24	Month-to-Month

996	1997	Male	42	9	One-Year
997	1998	Female	29	3	Two-Year
998	1999	Male	18	31	Month-to-Month
999	2000	Male	39	40	One-Year

	MonthlyCharges (USD)	TotalCharges (USD)	Churn
0	75.756917	2659.067786	Yes
1	99.780443	2604.269557	No
2	91.728821	1238.339084	Yes
3	66.079873	2557.291102	Yes
4	93.391649	672.419873	Yes
..
995	70.236004	1517.097691	No
996	73.944305	598.948867	Yes
997	62.598495	169.015938	No
998	77.713765	2168.214054	No
999	31.302244	1126.880782	No

[1000 rows x 8 columns]

```
[13]: # gender to numeric/reference numbers
gender = pd.get_dummies(data=df['Gender'], drop_first=True)

# contractType to numeric/reference numbers
contractType = pd.get_dummies(data=df['ContractType'])

# Churn to numeric/reference numbers
churn = pd.get_dummies(df['Churn'], drop_first=True)
churn.columns=['Churn_Status']
```

```
[14]: df = pd.concat([df, churn, gender, contractType], axis=1)
df.head()
```

```
[14]: CustomerID  Gender  Age  ServiceLength (months)  ContractType \
0          1001   Male   66              39      Two-Year
1          1002  Female   27              29  Month-to-Month
2          1003   Male   45              15      Two-Year
3          1004   Male   38              43      Two-Year
4          1005   Male   67              8      Two-Year

MonthlyCharges (USD)  TotalCharges (USD)  Churn  Churn_Status  Male \
0          75.756917      2659.067786   Yes           1         1
1          99.780443      2604.269557   No           0         0
2          91.728821      1238.339084   Yes           1         1
3          66.079873      2557.291102   Yes           1         1
4          93.391649        672.419873   Yes           1         1
```

	Month-to-Month	One-Year	Two-Year
0	0	0	1
1	1	0	0
2	0	0	1
3	0	0	1
4	0	0	1

```
[15]: # delete string/unnecessary columns as they are not needed
df.drop(columns=['CustomerID', 'Gender', 'ContractType', 'Churn'], inplace=True)
```

```
[16]: df.head()
```

```
[16]:   Age  ServiceLength (months)  MonthlyCharges (USD)  TotalCharges (USD) \
0    66                    39          75.756917         2659.067786
1    27                    29          99.780443         2604.269557
2    45                    15          91.728821         1238.339084
3    38                    43          66.079873         2557.291102
4    67                     8          93.391649          672.419873
```

	Churn_Status	Male	Month-to-Month	One-Year	Two-Year
0	1	1	0	0	1
1	0	0	1	0	0
2	1	1	0	0	1
3	1	1	0	0	1
4	1	1	0	0	1

5 Data Splitting

```
[17]: # INDEPENDENT/FEATURES and LABELS/DEPENDENT VARIABLES/COLS
X = df.drop('Churn_Status', axis=1)
y = df['Churn_Status']
```

```
[18]: # SPLITTING DATA
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
                                                    random_state=42)
```

6 Logistic Regression

```
[19]: # CREATE MODEL AND TRAIN
from sklearn.linear_model import LogisticRegression

log_model = LogisticRegression()
log_model.fit(X_train, y_train)
```

```
[19]: LogisticRegression()
```

```
[20]: # PREDICT
      predictions = log_model.predict(X_test)
```

```
[21]: # EVALUATION
      from sklearn.metrics import classification_report, confusion_matrix

      print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.61	0.99	0.76	153
1	0.33	0.01	0.02	97
accuracy			0.61	250
macro avg	0.47	0.50	0.39	250
weighted avg	0.50	0.61	0.47	250

```
[22]: confusion_matrix(y_test, predictions)
```

```
[22]: array([[151,  2],
        [ 96,  1]], dtype=int64)
```

7 Random Forest

```
[23]: from sklearn.ensemble import RandomForestClassifier

      rf_model = RandomForestClassifier()
      rf_model.fit(X, y)
```

```
[23]: RandomForestClassifier()
```

```
[24]: predictions = rf_model.predict(X_test)
```

```
[25]: confusion_matrix(y_test, predictions)
```

```
[25]: array([[153,  0],
        [  0, 97]], dtype=int64)
```

```
[26]: print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	153
1	1.00	1.00	1.00	97

accuracy			1.00	250
macro avg	1.00	1.00	1.00	250
weighted avg	1.00	1.00	1.00	250

8 TESTING MODELS

```
[27]: test_data = {
    'Age': 45,
    'ServiceLength (months)': 30,
    'MonthlyCharges (USD)': 60,
    'TotalCharges (USD)': 30 * 60,
    'Male': 0,
    'Month-to-Month': 0,
    'One-Year': 1,
    'Two-Year': 0
}

test_df = pd.DataFrame([test_data])
test_df
```

```
[27]:   Age  ServiceLength (months)  MonthlyCharges (USD)  TotalCharges (USD) \
0   45                      30                      60             1800

   Male  Month-to-Month  One-Year  Two-Year
0     0               0         1         0
```

```
[28]: def check_churn_status(val):
    if val==0:
        return ("Customers is NOT expected to churn")
    elif val==1:
        return ("Customers is expected to churn")
```

```
[29]: # LOGISTIC REGRESSION TEST
churn_prediction_lr_model = log_model.predict(test_df)

print("Logistic Model Result\n")
check_churn_status(churn_prediction_lr_model[0])
```

Logistic Model Result

```
[29]: 'Customers is NOT expected to churn'
```

```
[30]: # RANDOM FOREST CALSSIFICATION TEST
churn_prediction_rf_model = rf_model.predict(test_df)

print("RANDOM FOREST CALSSIFICATION Model Result\n")
```

```
check_churn_status(churn_prediction_rf_model[0])
```

RANDOM FOREST CALSSIFICATION Model Result

```
[30]: 'Customers is NOT expected to churn'
```

9 FLASK API

```
[31]: from flask import Flask, request, jsonify

app = Flask(__name__)

@app.route('/')
def hello():
    return "Welcome to Churn Prediction API!"

@app.route('/predict', methods=['POST'])
def predict():
    try:

        test_data = {
            'Age': request.json['Age'],
            'ServiceLength (months)': request.json['ServiceLength (months)'],
            'MonthlyCharges (USD)': request.json['MonthlyCharges (USD)'],
            'TotalCharges (USD)': request.json['ServiceLength (months)] * ↵
↪request.json['MonthlyCharges (USD)'],
            'Male': request.json['Male'],
            'Month-to-Month': request.json['Month-to-Month'],
            'One-Year': request.json['One-Year'],
            'Two-Year': request.json['Two-Year']
        }

        test_df = pd.DataFrame([test_data])

        # Logistic Regression prediction
        churn_prediction_lr = log_model.predict(test_df)

        # Random Forest prediction
        churn_prediction_rf = rf_model.predict(test_df)

        return jsonify({
            "Logistic Regression Prediction": ↵
↪check_churn_status(churn_prediction_lr[0]),
            "Random Forest Prediction": ↵
↪check_churn_status(churn_prediction_rf[0])
        })
```

```

except Exception as e:
    return jsonify({"error": str(e)})

if __name__ == '__main__':
    app.run(host='0.0.0.0', port=5000)

```

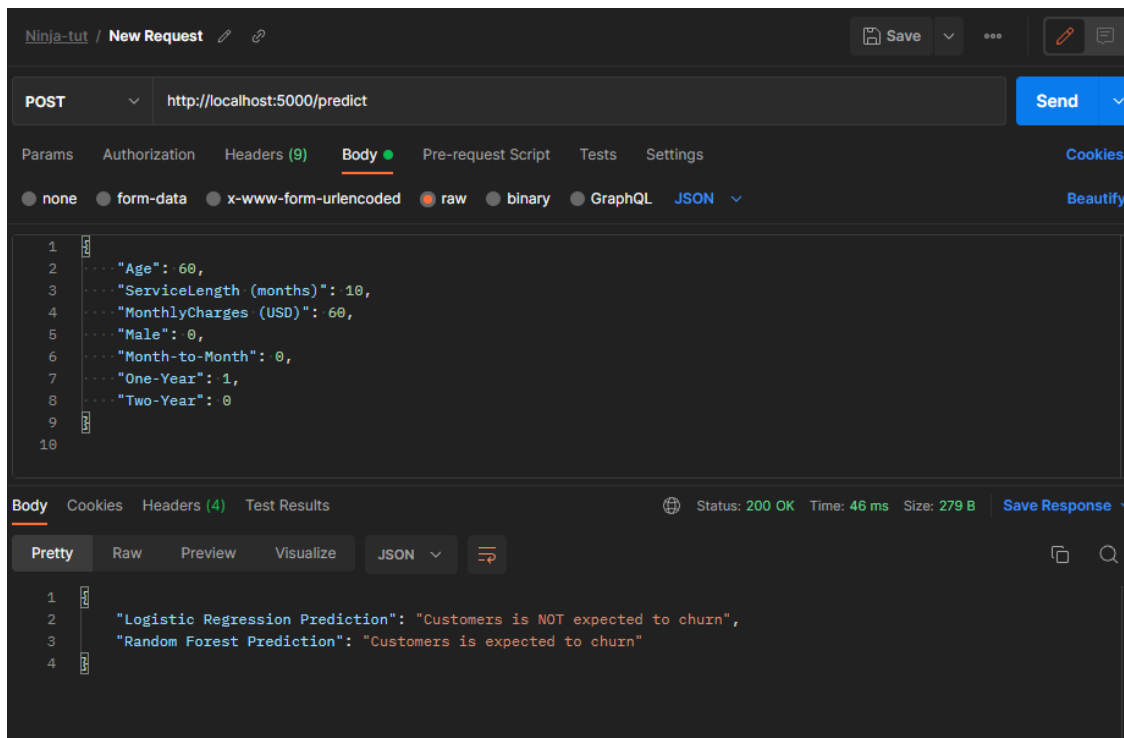
```

* Serving Flask app "__main__" (lazy loading)
* Environment: production
  WARNING: This is a development server. Do not use it in a production
deployment.
  Use a production WSGI server instead.
* Debug mode: off

* Running on all addresses.
  WARNING: This is a development server. Do not use it in a production
deployment.
* Running on http://192.168.100.207:5000/ (Press CTRL+C to quit)
127.0.0.1 - - [22/Aug/2023 18:32:12] "POST /predict HTTP/1.1" 200 -

```

10 Postman Output



[]: